

Prodigy Infotech Internship Task by Chaitanya gadekar

Task 3 : Decision Tree Classification on Bank Marketing Dataset

Problem Statement

Task-03



Build a decision tree classifier to predict whether a customer will purchase a product or service based on their demographic and behavioral data. Use a dataset such as the Bank Marketing dataset from the UCI Machine Learning Repository.

About the Dataset :

The Dataset we are going to use in this task is Bank Marketing Dataset which is taken from UCI Machine Learning Repository. The dataset here is the 10% sample of the Original Bank Marketing Dataset. The dataset contains 20 features and 1 label. The data is related with direct marketing campaigns of a Portuguese banking institution. The marketing campaigns were based on phone calls. Often, more than one contact to the same client was required, in order to access if the product (bank term deposit) would be ('yes') or not ('no') subscribed.

Features of the Dataset:

Input variables:

bank client data:

1. age (numeric)
2. job : type of job (categorical: 'admin.', 'blue-collar', 'entrepreneur', 'housemaid', 'management', 'retired', 'self-employed', 'services', 'student', 'technician', 'unemployed', 'unknown')

3. marital : marital status (categorical: 'divorced','married','single','unknown'; note: 'divorced' means divorced or widowed)
4. education (categorical: 'basic.4y','basic.6y','basic.9y','high.school','illiterate','professional.course','university.degree','unknown')
5. default: has credit in default? (categorical: 'no','yes','unknown')
6. housing: has housing loan? (categorical: 'no','yes','unknown')
7. loan: has personal loan? (categorical: 'no','yes','unknown') ## related with the last contact of the current campaign:
8. contact: contact communication type (categorical: 'cellular','telephone')
9. month: last contact month of year (categorical: 'jan', 'feb', 'mar', ..., 'nov', 'dec')
10. day_of_week: last contact day of the week (categorical: 'mon','tue','wed','thu','fri')
11. duration: last contact duration, in seconds (numeric). Important note: this attribute highly affects the output target (e.g., if duration=0 then y='no'). Yet, the duration is not known before a call is performed. Also, after the end of the call y is obviously known. Thus, this input should only be included for benchmark purposes and should be discarded if the intention is to have a realistic predictive model. ## other attributes:
12. campaign: number of contacts performed during this campaign and for this client (numeric, includes last contact)
13. pdays: number of days that passed by after the client was last contacted from a previous campaign (numeric; 999 means client was not previously contacted)
14. previous: number of contacts performed before this campaign and for this client (numeric)
15. poutcome: outcome of the previous marketing campaign (categorical: 'failure','nonexistent','success') ## social and economic context attributes
16. emp.var.rate: employment variation rate - quarterly indicator (numeric)
17. cons.price.idx: consumer price index - monthly indicator (numeric)
18. cons.conf.idx: consumer confidence index - monthly indicator (numeric)
19. euribor3m: euribor 3 month rate - daily indicator (numeric)
20. nr.employed: number of employees - quarterly indicator (numeric)

Output variable (desired target):

1. y - has the client subscribed a term deposit? (binary: 'yes','no')

The classification goal is to predict if the client will subscribe (yes/no) a term deposit (variable y).

Importing Libraries

```
import pandas as pd
# pandas is aliased as pd
import numpy as np
# numpy is aliased as np
import matplotlib.pyplot as plt
```

```
# pyplot is aliased as plt
import seaborn as sns
# seaborn is aliased as sns

import warnings
warnings.filterwarnings('ignore')

%matplotlib inline
```

Loading the Dataset

```
df = pd.read_csv('bank-additional.csv', delimiter=';')
df.rename(columns={'y': 'deposit'}, inplace=True)
df.head()
```

	age	job	marital	education	default	housing
loan \						
0	30	blue-collar	married	basic.9y	no	yes
no						
1	39	services	single	high.school	no	no
no						
2	25	services	married	high.school	no	yes
no						
3	38	services	married	basic.9y	no	unknown
unknown						
4	47	admin.	married	university.degree	no	yes
no						

	contact	month	day_of_week	...	campaign	pdays	previous
poutcome \							
0	cellular	may	fri	...	2	999	0
nonexistent							
1	telephone	may	fri	...	4	999	0
nonexistent							
2	telephone	jun	wed	...	1	999	0
nonexistent							
3	telephone	jun	fri	...	3	999	0
nonexistent							
4	cellular	nov	mon	...	1	999	0
nonexistent							

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed
deposit					
0	-1.8	92.893	-46.2	1.313	5099.1
no					
1	1.1	93.994	-36.4	4.855	5191.0
no					
2	1.4	94.465	-41.8	4.962	5228.1

```
no
3      1.4      94.465      -41.8      4.959      5228.1
no
4     -0.1      93.200      -42.0      4.191      5195.8
no
```

```
[5 rows x 21 columns]
```

```
# Showing first 5 rows
df.head()
```

```
   age  job  marital  education  default  housing
loan \
0   30 blue-collar  married    basic.9y      no      yes
no
1   39  services   single    high.school      no      no
no
2   25  services   married    high.school      no      yes
no
3   38  services   married    basic.9y      no unknown
unknown
4   47   admin.   married  university.degree      no      yes
no
```

```
   contact month day_of_week  ...  campaign  pdays  previous
poutcome \
0   cellular    may         fri  ...        2    999          0
nonexistent
1   telephone  may         fri  ...        4    999          0
nonexistent
2   telephone  jun         wed  ...        1    999          0
nonexistent
3   telephone  jun         fri  ...        3    999          0
nonexistent
4   cellular   nov         mon  ...        1    999          0
nonexistent
```

```
   emp.var.rate  cons.price.idx  cons.conf.idx  euribor3m  nr.employed
deposit
0          -1.8          92.893          -46.2        1.313        5099.1
no
1          1.1          93.994          -36.4        4.855        5191.0
no
2          1.4          94.465          -41.8        4.962        5228.1
no
3          1.4          94.465          -41.8        4.959        5228.1
no
4         -0.1          93.200          -42.0        4.191        5195.8
no
```

```
[5 rows x 21 columns]
```

```
# showing last 5 rows
```

```
df.tail()
```

	age	job	marital	education	default	housing	loan
contact \							
4114	30	admin.	married	basic.6y	no	yes	yes
cellular							
4115	39	admin.	married	high.school	no	yes	no
telephone							
4116	27	student	single	high.school	no	no	no
cellular							
4117	58	admin.	married	high.school	no	no	no
cellular							
4118	34	management	single	high.school	no	yes	no
cellular							

	month	day_of_week	...	campaign	pdays	previous	poutcome	\
4114	jul	thu	...	1	999	0	nonexistent	
4115	jul	fri	...	1	999	0	nonexistent	
4116	may	mon	...	2	999	1	failure	
4117	aug	fri	...	1	999	0	nonexistent	
4118	nov	wed	...	1	999	0	nonexistent	

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m
nr.employed \				
4114	1.4	93.918	-42.7	4.958
5228.1				
4115	1.4	93.918	-42.7	4.959
5228.1				
4116	-1.8	92.893	-46.2	1.354
5099.1				
4117	1.4	93.444	-36.1	4.966
5228.1				
4118	-0.1	93.200	-42.0	4.120
5195.8				

	deposit
4114	no
4115	no
4116	no
4117	no
4118	no

```
[5 rows x 21 columns]
```

Basic Understanding of the Dataset

```
# showing dimensions of the dataset
df.shape

(4119, 21)
```

The dataset contains 4119 rows and 21 columns

```
# showing column names
df.columns

Index(['age', 'job', 'marital', 'education', 'default', 'housing',
       'loan',
       'contact', 'month', 'day_of_week', 'duration', 'campaign',
       'pdays',
       'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',
       'cons.conf.idx', 'euribor3m', 'nr.employed', 'deposit'],
      dtype='object')

# checking for data types
df.dtypes

# checking for different data types
df.dtypes.value_counts()
```

```

float64      5
dtype: int64

# showing information about the dataset
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4119 entries, 0 to 4118
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   4119 non-null   int64
 1   job                   4119 non-null   object
 2   marital               4119 non-null   object
 3   education              4119 non-null   object
 4   default               4119 non-null   object
 5   housing               4119 non-null   object
 6   loan                  4119 non-null   object
 7   contact               4119 non-null   object
 8   month                 4119 non-null   object
 9   day_of_week           4119 non-null   object
10   duration              4119 non-null   int64
11   campaign              4119 non-null   int64
12   pdays                 4119 non-null   int64
13   previous              4119 non-null   int64
14   poutcome              4119 non-null   object
15   emp.var.rate          4119 non-null   float64
16   cons.price.idx         4119 non-null   float64
17   cons.conf.idx         4119 non-null   float64
18   euribor3m             4119 non-null   float64
19   nr.employed           4119 non-null   float64
20   deposit               4119 non-null   object
dtypes: float64(5), int64(5), object(11)
memory usage: 675.9+ KB

```

From the above information we can conclude that -

1. The dataset has 21 columns and 4119 rows.
2. The dataset has 11 categorical columns
3. The dataset has 10 numerical columns
4. the dataset has no null values

Data Cleaning and Data Preprocessing

Handling Duplicated Values

```
# checking for duplicates
df.duplicated().sum()

0
```

Handling Null Values

```
df.isna().sum()
```

There is no null values in the dataset

Extracting Numerical and Categorical Columns

```
cat_cols = df.select_dtypes(include='object').columns
print(cat_cols)

num_cols = df.select_dtypes(exclude='object').columns
print(num_cols)

Index(['job', 'marital', 'education', 'default', 'housing', 'loan',
      'contact',
      'month', 'day_of_week', 'poutcome', 'deposit'],
```



```

dtype='object')
Index(['age', 'duration', 'campaign', 'pdays', 'previous',
      'emp.var.rate',
      'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed'],
      dtype='object')

```

Descriptive Statistical Analysis

```

# For Numerical Columns
df.describe()

```

	age	duration	campaign	pdays	previous
\					
count	4119.000000	4119.000000	4119.000000	4119.000000	4119.000000
mean	40.113620	256.788055	2.537266	960.422190	0.190337
std	10.313362	254.703736	2.568159	191.922786	0.541788
min	18.000000	0.000000	1.000000	0.000000	0.000000
25%	32.000000	103.000000	1.000000	999.000000	0.000000
50%	38.000000	181.000000	2.000000	999.000000	0.000000
75%	47.000000	317.000000	3.000000	999.000000	0.000000
max	88.000000	3643.000000	35.000000	999.000000	6.000000

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m
nr.employed				
count	4119.000000	4119.000000	4119.000000	4119.000000
4119.000000				
mean	0.084972	93.579704	-40.499102	3.621356
5166.481695				
std	1.563114	0.579349	4.594578	1.733591
73.667904				
min	-3.400000	92.201000	-50.800000	0.635000
4963.600000				
25%	-1.800000	93.075000	-42.700000	1.334000
5099.100000				
50%	1.100000	93.749000	-41.800000	4.857000
5191.000000				
75%	1.400000	93.994000	-36.400000	4.961000
5228.100000				
max	1.400000	94.767000	-26.900000	5.045000
5228.100000				

```
# For Categorical columns
df.describe(include='object')
```

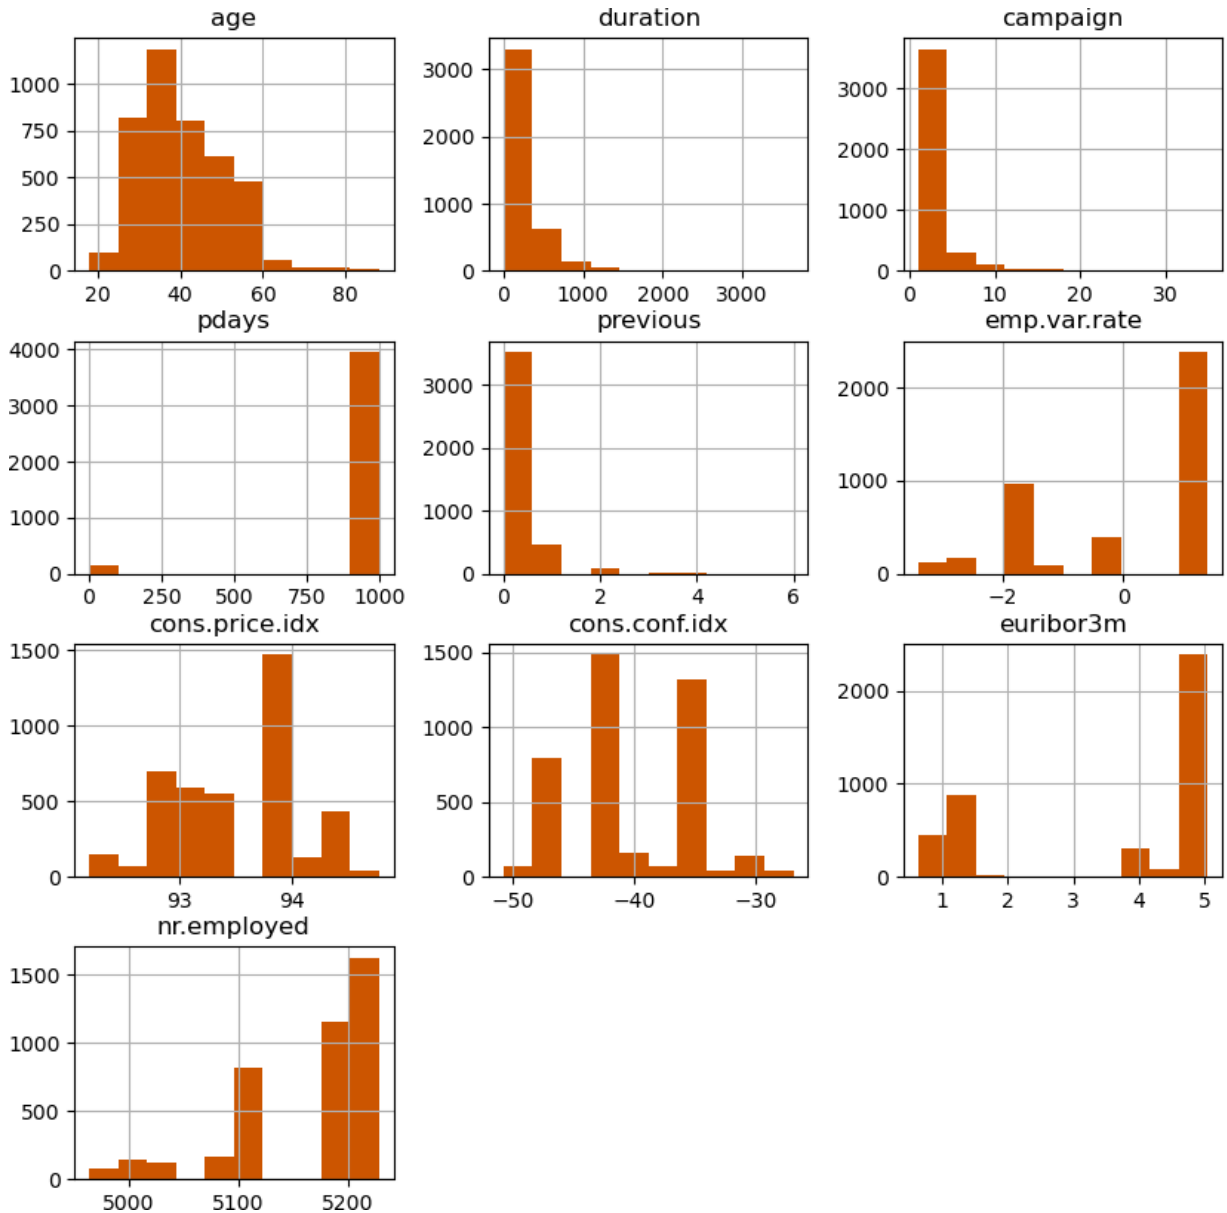
	job	marital	education	default	housing	loan
contact \						
count	4119	4119	4119	4119	4119	4119
unique	12	4	8	3	3	3
top	admin.	married	university.degree	no	yes	no
freq	1012	2509	1264	3315	2175	3349

	month	day_of_week	poutcome	deposit
count	4119	4119	4119	4119
unique	10	5	3	2
top	may	thu	nonexistent	no
freq	1378	860	3523	3668

Data Visualization


Visualizing Numerical columns using Histplot


```
df.hist(figsize=(10,10),color='#cc5500')
plt.show()
```





Visualizing Categorical columns using Barplot


```
for feature in cat_cols:
    plt.figure(figsize=(5,5)) # Adjust the figure size as needed
    sns.countplot(x=feature, data=df, palette='Wistia')
    plt.title(f'Bar Plot of {feature}')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.xticks(rotation=90)
    plt.show()
```
















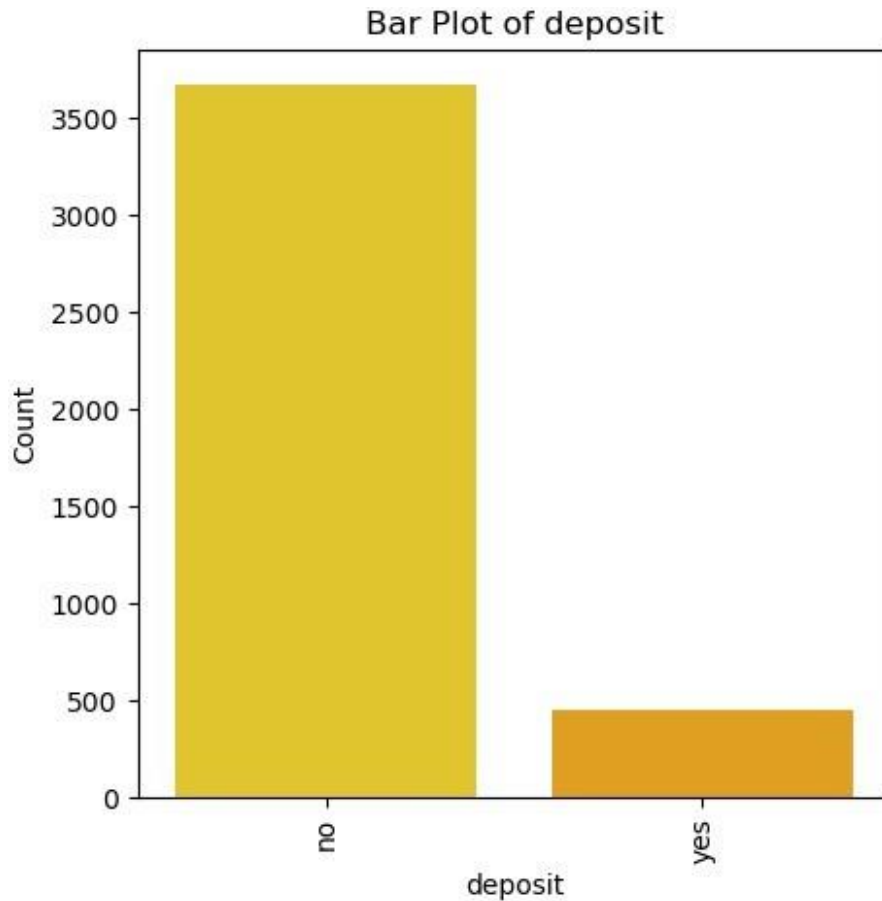










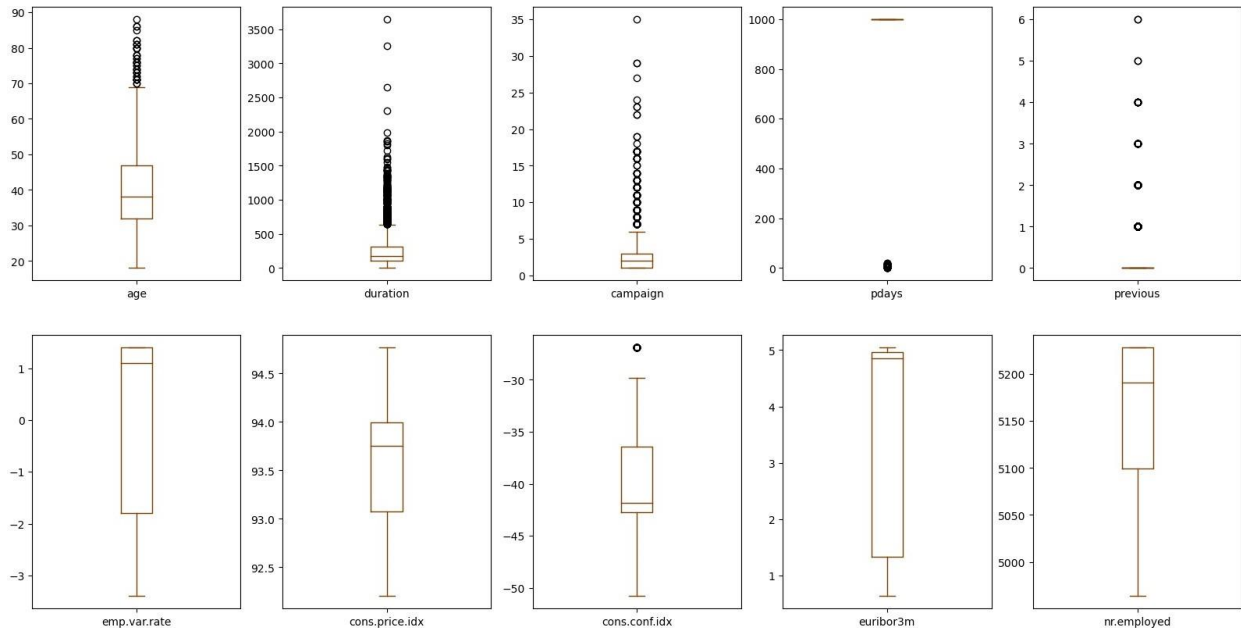


Insights:

1. In the Job Column, we have seen most of the clients are working as 'admin'.
2. In the marital Column, we have seen most of the clients are married.
3. In the education Column, we have seen most of the clients are having 'university.degree' as education.
4. In the default Column, we have seen most of the clients are having 'no' credit as default.
5. In the housing Column, we have seen most of the clients are taking housing loan.
6. In the loan Column, we have seen most of the clients are not taking personal loan.
7. In the contact Column, we have seen most of the clients are choosen cellular as contact.
8. In the month Column, we have seen most of the clients are contacted in the 'may' month.
9. In the day_of_week Column, we have seen most of the clients are contacted in 'thursday'.
10. In the poutcome Column, we have seen the result of most of the previous market campaign is 'nonexistent'.
11. In the target column , we have seen most of the clients are not subscribed a term deposit.

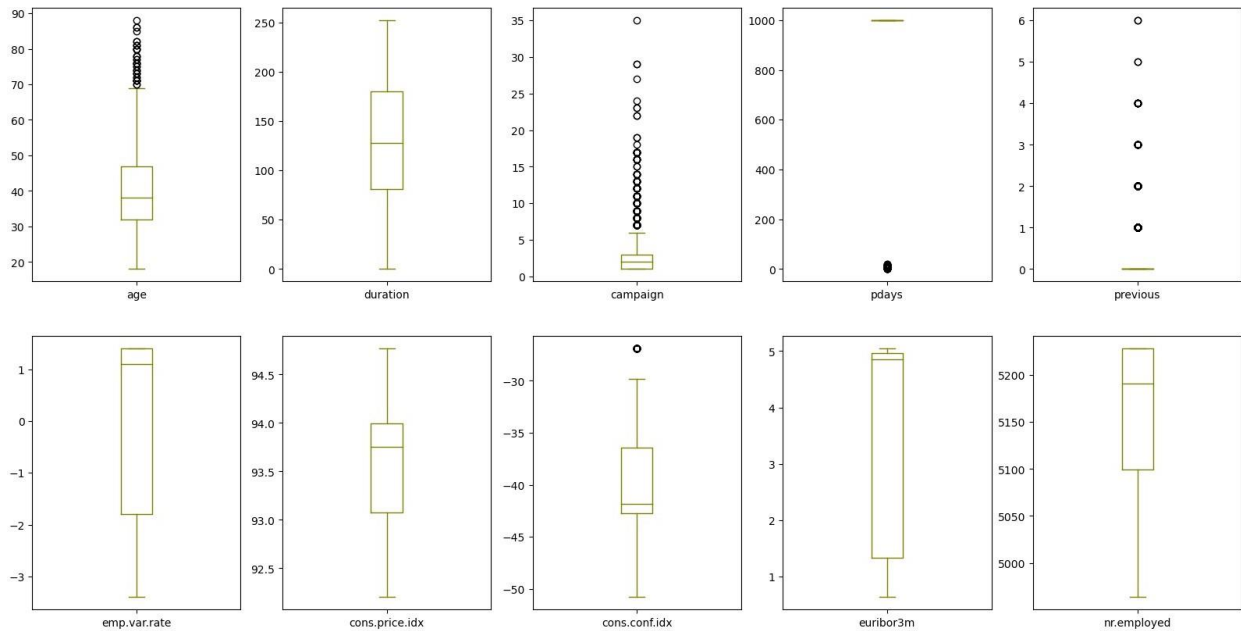
Plotting BoxPlot and Checking for Outliers

```
df.plot(kind='box', subplots=True,  
layout=(2,5),figsize=(20,10),color='#7b3f00')  
plt.show()
```



Through this Plot we can see there are 3 columns having outliers i.e.. 'Age', 'duration' and 'Campaign'. So, we will remove these outliers using Interquantile Range.

```
# Removing outliers  
  
column = df[['age', 'campaign', 'duration']]  
q1 = np.percentile(column, 25)  
q3 = np.percentile(column, 75)  
iqr = q3 - q1  
lower_bound = q1 - 1.5 * iqr  
upper_bound = q3 + 1.5 * iqr  
df[['age', 'campaign', 'duration']] = column[(column > lower_bound) &  
(column < upper_bound)]  
  
# Plotting boxplot after removing outliers  
df.plot(kind='box', subplots=True,  
layout=(2,5),figsize=(20,10),color='#808000')  
plt.show()
```



Checking for correlation using Correlation Plot


```
corr = df.corr()
print(corr)
corr = corr[abs(corr)>=0.90]
sns.heatmap(corr,annot=True,cmap='Set3',linewidths=0.2)
plt.show()
```

	age	duration	campaign	pdays	previous	\
age	1.000000	0.014048	-0.014169	-0.043425	0.050931	
duration	0.014048	1.000000	-0.218111	-0.093694	0.094206	
campaign	-0.014169	-0.218111	1.000000	0.058742	-0.091490	
pdays	-0.043425	-0.093694	0.058742	1.000000	-0.587941	
previous	0.050931	0.094206	-0.091490	-0.587941	1.000000	
emp.var.rate	-0.019192	-0.063870	0.176079	0.270684	-0.415238	
cons.price.idx	-0.000482	-0.013338	0.145021	0.058472	-0.164922	
cons.conf.idx	0.098135	0.045889	0.007882	-0.092090	-0.051420	
euribor3m	-0.015033	-0.067815	0.159435	0.301478	-0.458851	
nr.employed	-0.041936	-0.097339	0.161037	0.381983	-0.514853	

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m
\				
age	-0.019192	-0.000482	0.098135	-0.015033
duration	-0.063870	-0.013338	0.045889	-0.067815
campaign	0.176079	0.145021	0.007882	0.159435
pdays	0.270684	0.058472	-0.092090	0.301478
previous	-0.415238	-0.164922	-0.051420	-0.458851

emp.var.rate	1.000000	0.755155	0.195022	0.970308
cons.price.idx	0.755155	1.000000	0.045835	0.657159
cons.conf.idx	0.195022	0.045835	1.000000	0.276595
euribor3m	0.970308	0.657159	0.276595	1.000000
nr.employed	0.897173	0.472560	0.107054	0.942589

	nr.employed
age	-0.041936
duration	-0.097339
campaign	0.161037
pdays	0.381983
previous	-0.514853
emp.var.rate	0.897173
cons.price.idx	0.472560
cons.conf.idx	0.107054
euribor3m	0.942589
nr.employed	1.000000



Feature Selection using Correlation

```
high_corr_cols = ['emp.var.rate', 'euribor3m', 'nr.employed']

# copy with original dataframe
df1 = df.copy()
df1.columns

Index(['age', 'job', 'marital', 'education', 'default', 'housing',
      'loan',
      'contact', 'month', 'day_of_week', 'duration', 'campaign',
      'pdays',
      'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',
      'cons.conf.idx', 'euribor3m', 'nr.employed', 'deposit'],
      dtype='object')

# Removing high correlated columns from the dataset
df1.drop(high_corr_cols, inplace=True, axis=1) # axis=1 indicates
columns
df1.columns
```


	day_of_week	duration	campaign	pdays	previous	poutcome	\
0	0	250	1	20	0	1	
1	0	250	3	20	0	1	
2	4	224	0	20	0	1	
3	0	14	2	20	0	1	
4	1	55	0	20	0	1	
...	
4114	2	50	0	20	0	1	
4115	0	216	0	20	0	1	
4116	1	61	1	20	1	0	
4117	0	250	0	20	0	1	
4118	4	172	0	20	0	1	

	cons.price.idx	cons.conf.idx	deposit
0	8	4	0
1	18	16	0
2	23	8	0
3	23	8	0
4	11	7	0
...
4114	17	6	0
4115	17	6	0
4116	8	4	0
4117	13	17	0
4118	11	7	0

[4119 rows x 18 columns]

Checking for target variable

```
df_encoded['deposit'].value_counts()

0    3668
1     451
Name: deposit, dtype: int64
```

Selecting Independent and Dependent Variables

```
x = df_encoded.drop('deposit',axis=1) # independent variable
y = df_encoded['deposit']            # dependent variable
print(x.shape)
print(y.shape)
```

```
print(type(x))
print(type(y))

(4119, 17)
(4119,)
<class 'pandas.core.frame.DataFrame'>
<class 'pandas.core.series.Series'>
```

Splitting the dataset into Train and Test datasets

```
from sklearn.model_selection import train_test_split

print(4119*0.25)

1029.75

x_train,x_test,y_train,y_test =
train_test_split(x,y,test_size=0.25,random_state=1)
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)

(3089, 17)
(1030, 17)
(3089,)
(1030,)
```

Creating a function to compute Confusion Matrix, Classification Report and to generate training and testing scores

```
from sklearn.metrics import
confusion_matrix,classification_report,accuracy_score

def eval_model(y_test,y_pred):
    acc = accuracy_score(y_test,y_pred)
    print('Accuracy_Score',acc)
    cm = confusion_matrix(y_test,y_pred)
    print('Confusion Matrix\n',cm)
    print('Classification Report\n',
classification_report(y_test,y_pred))
```

```
def mscore(model):
    train_score = model.score(x_train,y_train)
    test_score = model.score(x_test,y_test)
    print('Training Score',train_score) # Training Accuracy
    print('Testing Score',test_score) # Testing Accuracy
```

Decision Tree Classifier

```
# Importing Decision Tree library
from sklearn.tree import DecisionTreeClassifier

# Building Decision Tree Classifier Model
dt =
DecisionTreeClassifier(criterion='gini',max_depth=5,min_samples_split=
10)
dt.fit(x_train,y_train)

DecisionTreeClassifier(max_depth=5, min_samples_split=10)

# Evaluating training and testing accuracy
mscore(dt)

Training Score 0.9148591777274199
Testing Score 0.8990291262135922

# Generating prediction
ypred_dt = dt.predict(x_test)
print(ypred_dt)

[0 0 1 ... 0 0 0]

# # Evaluate the model - confusion matrix, classification Report,
Accuracy
eval_model(y_test,ypred_dt)

Accuracy_Score 0.8990291262135922
Confusion Matrix
[[905  25]
 [ 79  21]]
Classification Report
```

	precision	recall	f1-score	support
0	0.92	0.97	0.95	930
1	0.46	0.21	0.29	100
accuracy			0.90	1030
macro avg	0.69	0.59	0.62	1030
weighted avg	0.87	0.90	0.88	1030


Plot Decision Tree

```
from sklearn.tree import plot_tree

# cn = class names, fn = feature_names
cn = ['no', 'yes']
fn = x_train.columns
print(fn)
print(cn)

Index(['age', 'job', 'marital', 'education', 'default', 'housing',
      'loan',
      'contact', 'month', 'day_of_week', 'duration', 'campaign',
      'pdays',
      'previous', 'poutcome', 'cons.price.idx', 'cons.conf.idx'],
      dtype='object')
['no', 'yes']

plot_tree(dt, feature_names=fn, class_names=cn, filled=True)
plt.show()
```



Decision Tree 2 (using entropy criteria) for visualization

```
# Building Decision Tree Classifier Model
dt1 =
DecisionTreeClassifier(criterion='entropy',max_depth=4,min_samples_split=15)
dt1.fit(x_train,y_train)

DecisionTreeClassifier(criterion='entropy', max_depth=4,
min_samples_split=15)

# Evaluating training and testing accuracy
mscore(dt1)

Training Score 0.9080608611201036
Testing Score 0.9048543689320389

# Generating prediction
ypred_dt1 = dt1.predict(x_test)

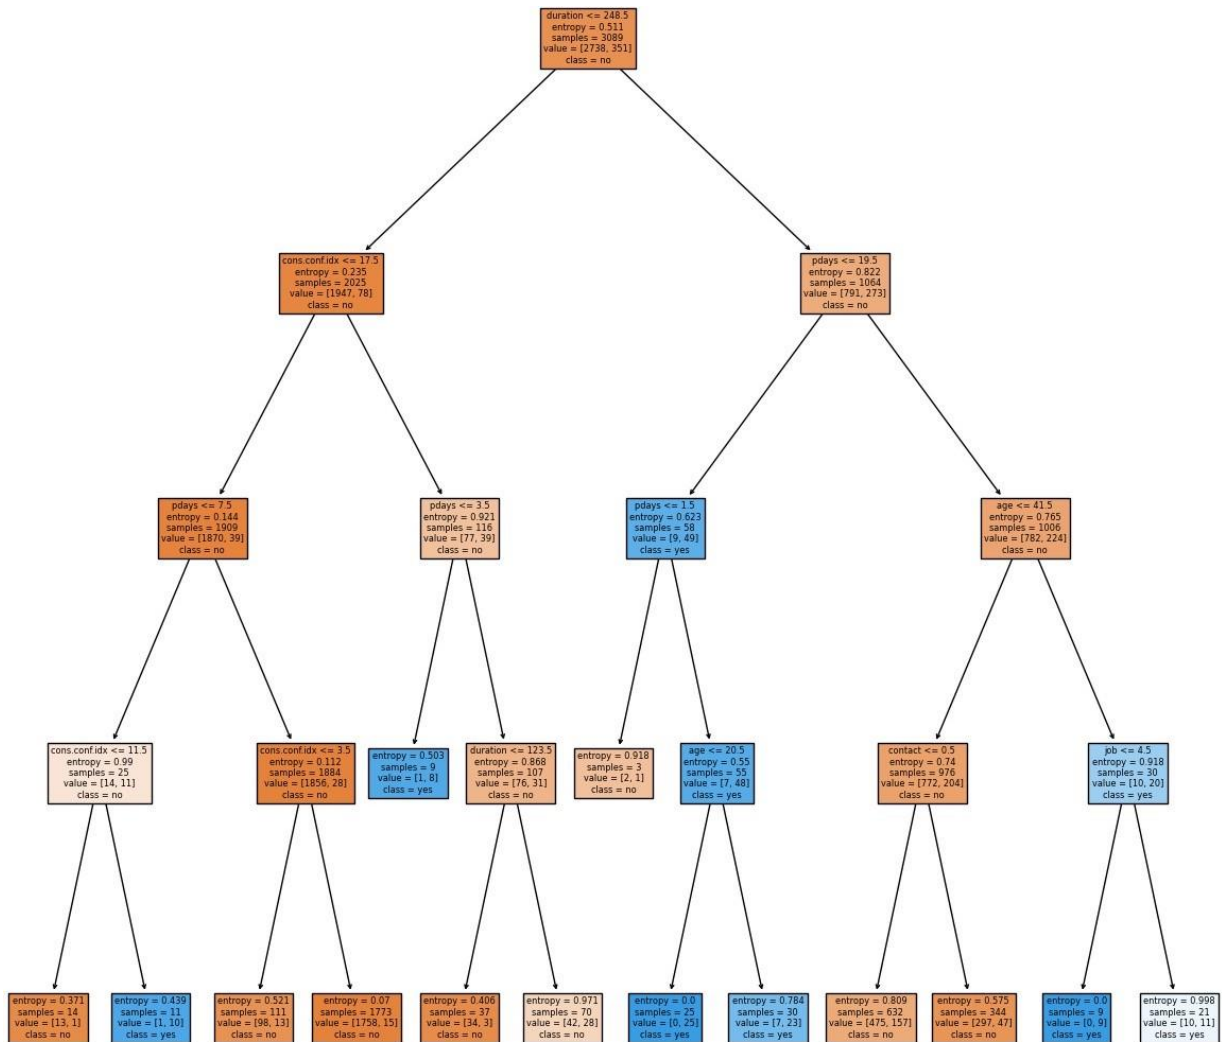
# Evaluate the model - confusion matrix, classification Report,
Accuracy
eval_model(y_test,ypred_dt1)

Accuracy_Score 0.9048543689320389
Confusion Matrix
[[915  15]
 [ 83  17]]
Classification Report

```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	930
1	0.53	0.17	0.26	100
accuracy			0.90	1030
macro avg	0.72	0.58	0.60	1030
weighted avg	0.88	0.90	0.88	1030

```
plt.figure(figsize=(15,15))
plot_tree(dt1,feature_names=fn,class_names=cn,filled=True)
plt.show()
```

Conclusion

In this we are having bank marketing dataset, which is sample dataset (10 %) of the original dataset. This dataset contains 4119 rows and 21 columns. In this there is one target variable which is deposit which implies whether the client has taken term deposit or not. The classification goal of this task is to predict if the client will subscribe (yes/no) a term deposit (variable y). In this after performing data cleaning, data preprocessing and feature selection, I have built the decision tree classifier model and evaluated the accuracy of the model which is 90 % using both the criterion (gini) and entropy. This accuracy score is high and the model is good for this dataset.

